

Extended cognition

Mark Sprevak
University of Edinburgh

10 February 2019

Article Summary

Extended cognition takes the idea that your mind is ‘on’ your smartphone literally. It says that human cognitive states and processes sometimes spill outside our heads and into objects in our environment. Alleged examples include not just smartphones, but also use of simpler technology (pencil and paper to perform a calculation), our own body (ticking off our fingers when we count), and other people (our spouse who remembers appointments so we don’t have to).

There are three main arguments for extended cognition.

Functionalism arguments rely on similarities in functional structure between extended processes and (actual or possible) internal cognitive processes. Cognition extends because the physical mechanisms that support it ‘work the same way’ in both cases. *Inference to the best explanation* arguments canvas the benefits that extended cognition would bring to psychology. We should believe that cognition extends because it would make our psychological theories more unified, elegant, and fruitful. *Second-wave* arguments emphasise the ways in which our brains integrate with our environment. Cognition extends because brains, bodies, and environment are so tightly intertwined that, when we solve certain cognitive tasks, they count as a single system.

Extended cognition is attacked on many fronts. It has been claimed that it generates *absurdly high levels of extension* (‘cognitive bloat’); that it is inferior to the more conservative hypothesis of *embedded cognition*; that its arguments confuse *causal*

coupling with constitution; and that its alleged cases fail to satisfy some *mark of the cognitive*.

Extended cognition concerns only the cognitive, information-processing aspects of mental life. It has, however, inspired similar claims about extension for other aspects of the mind, including conscious experience, emotions, moods, intentional agency, knowledge, and selfhood.

Extended cognition is part of a wider ‘4E cognition’ research programme. The four Es stand for extended, embedded, embodied, and enactive cognition. Each E offers a closely related, albeit distinct perspective on the role of the environment in cognition. Other forms of externalism about the mind – [content externalism](#), [direct realism about perception](#), [collective intentionality](#), and [group cognition](#) – are less closely related to extended cognition.

1 Extended cognition

Humans love their devices. Smartphones, tablets, computers, paper notebooks, calendars, and to-do lists play a pervasive role in our lives. These pieces of technology cause us to think and to behave in certain ways, and we in turn modify them to influence our future thought and behaviour. Often we don’t appreciate how much we rely on them until they’re gone. Empirical work in psychology suggests that this dependence runs very deep. Even during relatively undemanding tasks – e.g. copying a simple coloured pattern made of puzzle pieces – we off-load information processing onto the environment to reduce work for our brains (Ballard et al., 1997). Once one recognises this, one sees it everywhere: a bartender lines up cocktail glasses of different shapes to remember a complex order; a mathematician uses pencil and paper to guide her steps in a calculation; a child uses her fingers to count off the days until her next birthday. These observations reveal that intelligent, adaptive human thought and behaviour need not always be produced by the brain alone. It often involves a two-way cooperative interaction between the brain, body, and world (Dennett, 1996; Hutchins, 1995; Simon, 1969).

The hypothesis of extended cognition (HEC) goes beyond this relatively uncontroversial observation in a controversial way. Environmental processes don’t merely interact with our brain processes to produce thought and behaviour. Those environmental processes *have as much claim to be mental or cognitive* as their brain-based collaborators. Human cognition literally extends into smartphones, tablets, notebooks, to-do lists, and cocktail glasses.

HEC Some actual human cognitive states/processes extend outside our brains to include, as parts, states/processes in the environment.

HEC is frequently misunderstood, so it is worth taking care here.

HEC applies the slightly puzzling concept of ‘extension’ to cognitive processes. It is important to realise that HEC is not a claim about cognition considered in some abstract otherworldly sense, or about cognition considered from the first-person phenomenological point of view. It is a claim about cognition’s physical, mechanistic basis. That physical basis has a location and an extent. Traditional wisdom has it that the physical basis of human cognition (its ‘realiser’ in functionalist terms) is located exclusively inside the human brain. HEC says that this is false: a mix of brain and environmental states/processes sometimes underlie (realise) cognition.

What counts as the ‘environment’ may not always be precisely defined. For our purposes, the environment is anything outside the brain or central nervous system. Thus, your cognitive processes could extend into external technology (smartphones, notebooks), your non-neural body (fingers, limbs, tendons), naturally occurring objects around you (useful sticks and stones), tokens of public language (appropriate sound waves and ink marks), or the brains of trusted people (your spouse).

There are several claims that are often mistakenly associated with HEC.

HEC does not say that the environment is as ‘important’ as the brain for cognition. The brain may be the principal player behind our thought and behaviour. What HEC says is that it is not the *only* player: it is not the sole physical locus of human cognitive states/processes (see Clark, 2007a, on ‘organism-centered’ cognition). HEC does not say ‘how much’ of cognition extends. It is compatible with some, even most, of cognition being wholly internal and brain-based. Only *some* instances of human cognition need extend for HEC to be true. HEC is also silent about whether extended cognition is a ‘good thing’ for the agent in question. Sometimes it is (e.g. when a smartphone increases your memory capacity or improves your calculating ability), but sometimes it isn’t (e.g. when a smartphone sustains unhealthy cognitive biases or pathological worries about self image, see Sprevak, 2011; Borsboom, Cramer and Kalis, 2018). HEC similarly says nothing about whether extended cognition is the outcome of a deliberate choice by the agent (or anyone else). We may be drawn into entanglements with the environment without our awareness or consent. Finally, HEC does not say anything about whether objects like smartphones have an independent mental life of their own. As far as HEC is concerned, a smartphone sitting by itself need have no more mental life than an isolated neuron sitting in a Petri dish. It is only when a smartphone (or a neuron) stands in the right relationship to the rest of our brain that it becomes part of our cognitive life and thereby endowed with mental/cognitive properties.

2 Arguments for extended cognition

There are three main arguments for the hypothesis of extended cognition (HEC). All of them appear in Clark and Chalmers' (1998) original paper and each can be illustrated using their (now classic) thought experiment:

Inga has normal human memory. One day, she hears of an interesting exhibition at the Museum of Modern Art (MoMA). Inga thinks, recalls that MoMA is on 53rd Street, and sets off. Otto suffers from a mild form of Alzheimer's disease and he always writes down useful information in his notebook. On the same day, Otto also hears of the exhibition at MoMA, retrieves the address from his notebook, and sets off.

Clark and Chalmers say that Otto has a belief that MoMA is on 53rd Street. This belief – and Otto's cognitive processes associated with storing and recalling it – extends outside his brain and into his notebook. Before proceeding to the arguments, it is worth noting that this is not meant to be a realistic description of Alzheimer's disease or any actual real-life case. Clark and Chalmers are describing an idealised set of interactions with an environment that *would* give rise to extended cognition. They claim that more complex, although essentially similar, versions of these interactions occur in common real-life cases (e.g. smartphone use).

Functionalism arguments. The idea here is that Otto and Inga are similar in ways that matter to the philosophical theory of functionalism. **Functionalism** says that the functional role of a physical state/process determines whether that state/process is mental/cognitive. A functionalist about beliefs says that what makes a brain state a belief is that it functions in a way appropriate to a belief. Clark and Chalmers observe that Otto's notebook functions for him at least roughly in the same way as Inga's biological memory functions for her. The state of Otto's notebook interacts with his desires and other beliefs in a similar manner to that in which the state of Inga's biological memory interacts with her desires and other beliefs. Exposure to new information causes Otto to modify the state of his notebook. Exposure to new information causes Inga to modify the state of her biological memory. Otto's notebook causes him to stop at 53rd Street. Inga's biological memory causes her to stop at 53rd Street. The 'coarse-grained' functional role of the stored information appears to be the same in both cases. Clark and Chalmers conclude that just as Inga has a belief that MoMA is on 53rd Street, so Otto has a belief (partly realised in his notebook) with the same content.

This argument relies on accepting functionalism about beliefs, and also on the 'parity principle':

If, as we confront some task, a part of the world functions as a process

which, *were it done in the head*, we would have no hesitation in recognizing as part of the cognitive process, then that part of the world *is* (so we claim) part of the cognitive process. (Clark and Chalmers, 1998, p. 8)

The parity principle says that we should not be prejudiced against treating Otto's notebook as realising one of his cognitive states merely because it is outside his head. Recent versions of the functionalist argument for HEC exploit the parity principle to block an obvious objection to the original functionalist argument: that Otto's notebook and Inga's biological memory are not *exactly* alike in their functional roles. This flaw was pressed by early critics of HEC (Adams and Aizawa, 2001; Rupert, 2004). Otto and Inga are not exactly functionally alike, but Otto *is* functionally identical to a hypothetical intelligent being who uses a notebook inside its head to store information (e.g. an extraterrestrial who uses internal ink-marks to encode some of its memories). If it is *possible* for such an intelligent being to have an internal memory system like Otto's notebook, and we accept the parity principle and functionalism, then we should say that Otto's *actual* notebook counts as a case of extended cognition. (For a discussion of this argument, see Sprevak, 2009)

Inference to the best explanation. The idea here is that HEC affords some *explanatory advantage* to psychology. The argument takes the form of an *inference to the best explanation* (IBE): we should believe that HEC is true because of the explanatory benefits its acceptance would bring to psychology. Clark and Chalmers say that HEC would unify otherwise unrelated patterns of human thought and behaviour. Otto and Inga illustrate this point. Inga arrived at 53rd Street because she *wanted* to go to MoMA and *believed* that it was there. If HEC is true, then Otto arrived at 53rd Street because he *wanted* to go to MoMA and *believed* that it was there. If we assume HEC, we can describe both agents as instantiating the same belief–desire decision-making psychology, irrespective of whether they use internal or external resources for their information storage. The alternative – explaining Otto's thought and behaviour always in terms of his internal beliefs *about* his notebook – seems, by contrast, to introduce unnecessary complexity. The notebook is a constant in Otto's life, similar to Inga's biological memory. At least within some explanatory contexts, it would seem redundant to point to it when explaining Otto's behaviour. Accepting HEC promises to make the notion of belief in psychology 'deeper and more unified, and ... more useful in explanation' (Clark and Chalmers, 1998, p. 14). (For a discussion of this argument, see Sprevak, 2010)

Second-wave arguments. These arguments tend to downplay the considerations behind the functionalist and IBE arguments and instead emphasise 'complementarity'. What is meant by this is that external states and processes need not be similar – either functionally or in terms of their role in psychological explanation – to any

internal ones. What justifies HEC is instead a tight, reciprocal *integration* between brains, bodies, and the environment. This approach has affinities with dynamical systems theory, which models the brain, body, and environment as a single system whose elements cannot be prised apart when solving a cognitive task (Thelen and Smith, 1994). In their original paper, Clark and Chalmers call this position ‘active externalism’: all parts (internal and external) are ‘active’ in the sense that they work together in a non-redundant fashion to solve a cognitive task. The resulting interaction need not be similar to any (actual or possible) internal cognitive process, and there is no reason why it should be classified as such by psychology. Otto’s notebook isn’t cognitive because it is functionally like Inga’s memory, or because it is explanatorily beneficial for psychology to group them together. Otto’s notebook is cognitive because it is systematically, reciprocally, and inextricably integrated with his brain during certain cognitive tasks. This integration is intended, at least in some cases, to also have a normative dimension: it arises from and it is maintained by social practices, norms, and institutions. (For a reconstruction of second-wave arguments, see Menary, 2010a; Sutton, 2010) (Note that talk of a ‘second wave’ may suggest that the ‘first wave’ was a failure, but it is by no means clear that either the functionalist or IBE argument cannot be made to work, or that all three arguments for HEC cannot be combined into a unified package.)

3 Objections to extended cognition

The hypothesis of extended cognition (HEC) has, unsurprisingly, encountered considerable resistance. Here are four common objections.

Cognitive bloat. Fans of HEC say that human cognition extends in certain cases and to a certain degree, but they usually don’t want to say there is *too much* extension. Too much extension would risk rendering our existing mental concepts pointless, absurd, or otherwise unfit for purpose. We could end up saying you already believe everything on the Internet or that you and I believe the same things just because we happen to live in the same environment. This is called the problem of ‘cognitive bloat’. Avoiding cognitive bloat requires drawing some principled line between those relationships with the environment (functional/explanatory/integrational) that entail cognitive extension and those that do not. This has not proved easy. Consider that not only your smartphone, but also your office computer, local library, friends, co-workers, and in many cases the whole Internet contain information on which you routinely draw to guide your behaviour in a way not dissimilar to that in which Otto relies on his notebook. Consider too that you and I often draw on the same set of environmental resources to guide our respective behaviours. An argument for HEC needs to guard against becoming a *reductio* of the view it is trying to defend. Introducing restrictions to avoid cognitive bloat, however, may

cut against assumptions required to motivate any form of HEC at all – for example, they may violate the parity principle (see Sprevak, 2009).

Embedded cognition. The hypothesis of embedded cognition (HEMC) is an alternative to HEC that claims to offer all of the latter's explanatory benefits but without the 'cost' of saying that notebooks have mental/cognitive properties. HEMC and HEC agree on almost all of the key points: extended processes are functionally similar to internal cognitive processes; grouping extended and internal processes together is sometimes useful in psychological explanation; brains are integrated with the environment when we solve certain cognitive problems. HEMC and HEC differ in that HEC says that extended cognitive processes are 100% mental/cognitive whereas HEMC says that only the internal, brain-based parts of those processes are mental/cognitive. According to HEMC, extended processes are composed of separate mental and non-mental parts – the mental parts residing exclusively inside the head. According to HEC, extended processes have a uniformly mental/cognitive nature, inside and out. Critics of HEC argue that we should abandon HEC for HEMC because the latter provides the benefits of the former but it is more ontologically parsimonious (it attributes fewer mental/cognitive properties to the world) and it is more conservative (it fits better with traditional internalist psychology) (see Rupert, 2004; Sprevak, 2010).

The coupling/constitution fallacy. Arguments for extended cognition often describe a two-way coupling relationship between our brains and the environment. One might be tempted to think that cognition extends *because* of this coupling relationship. However, just because X and Y are coupled does not mean that X is part of Y. The expansion of a bimetallic strip in a thermostat is coupled to the motion of atoms in the surrounding air in the room, but that does not mean that the strip's expansion is a process that extends into the atoms of the air (Adams and Aizawa, 2007, p. 91). Similarly, your brain and the environment may be coupled but that does not mean that your brain processes extend into the environment. This objection appears to most directly threaten second-wave arguments for extended cognition, since those arguments emphasise tight, two-way coupling between the brain, body, and environment. An advocate of HEC needs to identify something beyond mere coupling that justifies cognitive extension. A natural resort would be to appeal to the *explanatory value* of treating the coupled system as cognitive or the *functional similarity* between the coupled system and an internal (actual or possible) case of cognition. However, this would collapse 'second-wave' arguments into their 'first-wave' versions (see Adams and Aizawa, 2001; Adams and Aizawa, 2007). Alternatively, a second-wave defender might stress the normative nature of the coupling relationship, although it is not obvious how this renders the relationship constitutive.

The mark of the cognitive. A 'mark of the cognitive' is a set of conditions that are either

necessary or sufficient for a physical state/process to be cognitive. A mark of the cognitive is often suggested as a way of overcoming impasses between philosophers who disagree about HEC. If all sides could agree on a mark of the cognitive, they could use it to decide whether disputed cases (like Otto's notebook) really are cases of extended cognition. Proposed marks of the cognitive tend, however, to be at least as controversial as the cases they are supposed to help decide. Proposed marks include: (i) a cognitive process must involve non-derived representational content and be functionally similar to actual cases of internal, brain-based human cognition (Adams and Aizawa, 2007); (ii) a cognitive process must be part of an integrated, persisting system, and it must causally contribute to a wide range of cognitive phenomena (Rupert, 2009); (iii) a process is cognitive so long as it 'belongs' to a subject in the sense of causally contributing to the intentional content of that subject's personal-level states (Rowlands, 2010). Critics of HEC tend to favour 'conservative' marks of the cognitive like (i) or (ii), which rule out many cases of extended cognition. Advocates of HEC tend to favour 'liberal' marks of the cognitive like (iii), which allow in many cases of extended cognition. There is no consensus about which is the 'correct' mark of the cognitive. One might also worry that there is no such a thing as a mark of the cognitive. Just as there is no mark that determines which physical processes are *healthy*, *living*, or *cancerous*, so there might be no set of necessary or sufficient conditions that determines which processes are *cognitive* (Allen, 2017).

4 Extended cognition's kindred views

Philosophical work has focused not only on arguments for, and objections against, the hypothesis of extended cognition (HEC), but also on tracing HEC's connections to other externalist claims about the mind. Two types of view related to HEC stand out: claims about other aspects of the mind extending; and other forms of philosophical externalism about cognition.

Other aspects of the mind extending. Clark and Chalmers (1998) restricted their extended-mind claim to the cognitive, information-processing aspects of the mind (e.g. storage and retrieval of dispositional beliefs). They flagged conscious experience as an aspect of the mind that is likely to escape their arguments: conscious experience is likely to be entirely brain-based. Noë (2004) disagrees: using variants of the IBE and second-wave arguments of §2, he argues that the mechanisms underlying visual consciousness extend outside the brain and into the body and environment. Clark (2009) replies that while such extension is theoretically possible, it is unlikely to happen because our brain's causal interface with the world (our eyes and our hands) has a narrow bandwidth which (information-rich) conscious experience cannot cross.

Consciousness is not the only target of novel claims about mental extension. Carter, Gordon and Palermos (2016) say that emotions extend when the physical processes involved in their [cognitive-appraisal](#) component extend. Colombetti and Roberts (2015) argue that all kinds of affective mental states (emotions, moods, sentiments, and temperaments) extend if external objects are used to regulate them. Vierkant (2014) argues that intentional agency and [willpower](#) extend. Pritchard (2010) and Carter and Kallestrup (2016) explore extension claims for epistemic states like [knowledge](#). Clark (2007b) argues that selfhood extends. Essays in Anderson et al. (2018) explore extension claims about other aspects of the mind, including processes that differ significantly from the task-based, problem-solving examples favoured by Clark and Chalmers (e.g. friendship and imaginative engagement with a text).

Other forms of externalism about cognition. HEC is often described as part of a wider research programme of ‘4E cognition’. The other Es are *embedded*, *embodied*, and *enactive* cognition. We have already met *embedded* cognition with the hypothesis of embedded cognition (see §3). *Embodied* cognition says that cognition depends in some way on the physical nature of our bodies. Different versions of embodied cognition cash this out differently. In some, embodied cognition is used to express a body-based version of HEC: cognition depends on our body because our physical bodies partly realise our cognitive processes (e.g. when a child counts on her fingers, her mathematical cognising extends into the physical movement of her fingers). However, embodied cognition may also be used to express other sorts of claims. Lakoff and Johnson say that the specific shape of the human body influences the semantic content of our concepts, including abstract logical and mathematical concepts (Johnson, 2007). This is something about which HEC has almost nothing to say. *Enactive* cognition says that cognition is constituted by a looping interaction between perception and action that involves brain, body, and world. This has obvious affinity with second-wave arguments for HEC (§2), but versions of enactivism might differ in various ways from HEC about how they see this integration working. Some enactivists reject HEC’s typical framing in terms of representations and information processing. ‘Sensorimotor’ enactivists say that cognition is constituted by implicit, non-representational bodily skills (Noë, 2004). ‘Autopoietic’ enactivists say that cognition is constituted by biologically inspired relationships, such as autopoiesis, adaptivity, and sense-making (Di Paolo, 2005).

HEC has a more distant relationship to other kinds of philosophical externalism, such as content externalism, direct realism about perception, collective intentionality, and group cognition. [Content externalism](#) claims that the representational content of our cognitive states does not supervene on the physical state of our brains. HEC has almost nothing to say about this. Its claim is exclusively about the location of the physical vehicles, not about the factors that determine their representational content (Hurley, 2010). [Direct realism about perception](#) says that perception in-

volves a perceiver standing in a special relation to a real external object beyond her brain. HEC assumes almost nothing about whether this view about perception is correct. Indeed, it is frequently stated in intentionalist (as opposed to direct-realist) terms. [Collective intentionality](#) is the claim that groups sometimes share mental states, such as belief, knowledge, and intention. Closely related to this is the notion of [group cognition](#): a cognitive state or process that is properly attributed to a group of individuals rather than to the individuals that compose that group. What HEC has in common with these two ideas is that cognitive states and processes need not reside exclusively inside individual human heads. However, unlike these two ideas, HEC is a claim about individual human cognition, not about group cognition. HEC is compatible with, but need not assert the existence of, shared mental states or group minds. HEC may be combined with a variety of externalist claims about the mind, but it makes a separate, unique claim about how the mind spreads out into world.

Related Articles

- [Belief](#)
- [Collective intentionality](#)
- [Consciousness](#)
- [Content: Wide and narrow](#)
- [Emotions, nature of](#)
- [Functionalism](#)
- [Inference to the best explanation](#)
- [Perception](#)
- [The realization of knowledge](#)
- [The will](#)

Articles that might consider linking to this article

- [Belief](#)
- [Collective intentionality](#)
- [Content: Wide and narrow](#)
- [Functionalism](#)
- [Metaphysics of knowledge](#)

Acknowledgements

I would like to thank Ken Aizawa, Andy Clark, Mike Wheeler, and Rob Rupert for helpful comments on a draft of this entry.

Bibliography

- Adams, F. and K. Aizawa (2001). “The bounds of cognition”. In: *Philosophical Psychology* 14, pp. 43–64.
(One of the earliest, but still one of the best, critical responses to extended cognition. Adams and Aizawa developed their points in more detail in their later book.)
- (2007). *The Bounds of Cognition*. Oxford: Blackwell.
(Building on their earlier work, Adams and Aizawa offer a detailed criticism of extended cognition. This book contains one of the best articulations of the coupling/constitution fallacy – see §3.)
- Allen, C. (2017). “On (not) defining cognition”. In: *Synthese* 194, pp. 4233–4249.
(Allen argues that there is no such thing as a mark of the cognitive – see §3.)
- Anderson, M., D. Cairns, M. Sprevak and M. Wheeler, eds. (2018). *A History of Distributed Cognition*. 4 volumes. Edinburgh: Edinburgh University Press.
(An anthology that explores how various claims about extended cognition were supported, and suppressed, in the past.)
- Ballard, D. H., M. M. Hayhoe, P. Pook and R. Rao (1997). “Deictic codes for the embodiment of cognition”. In: *Behavioral and Brain Sciences* 20, pp. 723–767.
(An example of empirical work that demonstrates off-loading information onto the environment when we solve a cognitive task. Rather than store a puzzle pattern in our brain, we just remember where to look in the world to find it when we need it.)
- Borsboom, D., A. Cramer and A. Kalis (2018). “Brain disorders? Not really... Why network structures block reductionism in psychopathology research”. In: *Behavioral and Brain Sciences*, pp. 1–54. DOI: [10.1017/S0140525X17002266](https://doi.org/10.1017/S0140525X17002266).
(Certain mental disorders are not solely brain-based.)
- Brooks, R. A. (1991). “Intelligence without representation”. In: *Artificial Intelligence* 47, pp. 139–159.
(A robot can generate intelligent, adaptive behaviour by off-loading its information processing and representation onto the environment.)
- Carter, J. A., E. C. Gordon and S. O. Palermos (2016). “Extended emotion”. In: *Philosophical Psychology* 29, pp. 198–217.
(Human emotions extend when their cognitive-appraisal component extends.)
- Carter, J. A. and J. Kallestrup (2016). “Extended cognition and propositional memory”. In: *Philosophy and Phenomenological Research* 92, pp. 691–714.

- (Explores links between extended cognition and extended claims in the epistemology of memory.)
- Clark, A. (2007a). "Curing cognitive hiccups: A defense of the extended mind". In: *The Journal of Philosophy* 106, pp. 163–192.
(Extended cognition offers explanatory benefits to psychology – see §2. Clark also clarifies the role of the brain, in his hypothesis of 'organism-centered' cognition.)
- (2007b). "Soft selves and ecological control". In: *Distributed Cognition and the Will*. Ed. by D. Ross, D. Spurrett, H. Kincaid and G. L. Stephens. MIT Press, pp. 101–122.
(Selves extend.)
- (2008). *Supersizing the Mind*. Oxford: Oxford University Press.
(This book pulls together many years of Clark's work on extended cognition. There is a foreword by David Chalmers and an Appendix that reprints their 'Extended mind' paper.)
- (2009). "Spreading the joy? Why the machinery of consciousness is (probably) still in the head". In: *Mind* 118, pp. 963–993.
(Conscious experience does not extend.)
- Clark, A. and D. J. Chalmers (1998). "The extended mind". In: *Analysis* 58, pp. 7–19.
(This much-cited paper remains one of the clearest and best articulations of extended cognition and the arguments behind it.)
- Colombetti, G. and T. Roberts (2015). "Extending the extended mind: the case for extended affectivity". In: *Philosophical Studies* 172, pp. 1243–1263.
(Affective mental states, including emotions, moods, sentiments, and temperaments, extend.)
- Dennett, D. C. (1978). "Where am I?" In: *Brainstorms*. Cambridge, MA: MIT Press, pp. 310–323.
(Dennett describes fitting radio transmitters to a human's neurons. He considers the implications of cognitive extension for traditional concepts of mind, personhood, and moral responsibility.)
- (1996). *Kinds of Minds*. New York, NY: Basic Books.
(Dennett emphasises our reciprocal interaction with external resources like public language. He describes a hierarchy of different kinds of minds, including 'Gregorian' minds whose states/processes are partly shaped by their interactions with external tools.)

- Di Paolo, E. A. (2005). "Autopoiesis, adaptivity, teleology, agency". In: *Phenomenology and the Cognitive Sciences* 4, pp. 429–452.
(An overview of autopoietic enactivism – see §4.)
- Hurley, S. (2010). "Varieties of externalism". In: *The Extended Mind*. Ed. by R. Menary. Cambridge, MA: MIT Press, pp. 101–153.
(Examines the relationship between extended cognition and other forms of externalism in philosophy.)
- Hutchins, E. (1995). *Cognition in the Wild*. Cambridge, MA: MIT Press.
(An influential anthropological study of how humans off-load cognition when navigating at sea.)
- Johnson, M. (2007). *The Meaning of the Body*. Chicago, IL: University of Chicago Press.
(The semantic content of our abstract concepts is shaped by the particular nature of our physical bodies – see §4.)
- Menary, R. (2010a). "Cognitive integration and the extended mind". In: *The Extended Mind*. Ed. by R. Menary. Cambridge, MA: MIT Press.
(Along with John Sutton's chapter in the same volume, this defends a second-wave argument for extended cognition – see §2.)
- ed. (2010b). *The Extended Mind*. Cambridge, MA: MIT Press.
(A useful anthology containing several chapters that develop second-wave arguments for extended cognition as well as back-and-forth exchanges between Clark and his critics.)
- Noë, A. (2004). *Action in Perception*. Cambridge, MA: MIT Press.
(Conscious visual experience extends, albeit in a way that should be understood in terms of 'sensorimotor' enactivism rather than extended cognition – see §4.)
- Pritchard, D. (2010). "Cognitive ability and the extended cognition thesis". In: *Synthese* 175, pp. 133–151.
(Explores the relationship between extended cognition and ability claims in epistemology.)
- Rowlands, M. (2010). *The New Science of the Mind: From Extended Mind to Embodied Phenomenology*. Cambridge, MA: MIT Press.
(Rowlands argues for a mark of the cognitive that is favourable to extended cognition – see §3.)
- Rupert, R. D. (2004). "Challenges to the hypothesis of extended cognition". In: *The Journal of Philosophy* 101, pp. 389–428.
(Rupert argues for a rival to the hypothesis of extended cognition, the hypothesis of embedded cognition – see §3.)

- Rupert, R. D. (2009). *Cognitive Systems and the Extended Mind*. Oxford: Oxford University Press.
(Rupert argues for a mark of the cognitive based around conditions of persistence and integration – see §3.)
- Shapiro, L. (2011). *Embodied Cognition*. London: Routledge.
(An excellent textbook covering extended cognition and related views.)
- Simon, H. A. (1969). *The Sciences of the Artificial*. Cambridge, MA: MIT Press.
(Simon’s ‘parable of the ant’, described in Chapter 3 of his book, shows how complex, adaptive behaviour can be generated by a simple internal mechanism coupled to the environment.)
- Sprevak, M. (2009). “Extended cognition and functionalism”. In: *The Journal of Philosophy* 106, pp. 503–527.
(The functionalist argument for extended cognition succeeds, but entails cognitive bloat – see §3.)
- (2010). “Inference to the hypothesis of extended cognition”. In: *Studies in History and Philosophy of Science* 41, pp. 353–362.
(IBE arguments for extended cognition do not work in the context of the hypothesis of embedded cognition – see §§2, 3.)
- (2011). “Neural sufficiency, reductionism, and cognitive neuropsychiatry”. In: *Philosophy, Psychiatry and Psychology* 18, pp. 339–344.
(The physical mechanisms underlying psychiatric disorders may not be brain-based.)
- Sutton, J. (2010). “Exograms and interdisciplinarity: History, the extended mind, and the civilizing process”. In: *The Extended Mind*. Ed. by R. Menary. Cambridge, MA: MIT Press, pp. 189–225.
(Along with Richard Menary’s chapter in the same volume, this presents a second-wave argument for extended cognition – see §2.)
- Thelen, E. and L. B. Smith (1994). *A Dynamical Systems Approach to the Development of Cognition and Action*. Cambridge, MA: MIT Press.
(Defends the research programme of modelling the brain, body, and world as a single, integrated system.)
- Vierkant, T. (2014). “Mental muscles and the extended will”. In: *Topoi* 33, pp. 57–65.
(The will and intentional agency extend.)
- Wheeler, M. (2005). *Reconstructing the Cognitive World*. Cambridge, MA: MIT Press.
(Explores links between extended cognition and phenomenological ideas about embodiment from Heidegger.)