

Predictive coding II: The computational level

Mark Sprevak
University of Edinburgh

19 October 2021

1 Introduction

When we encounter a new computing device, we often try to describe its computational characteristics in terms of the task it faces: this shop's cash register has the task of *adding numbers*, this computer programme has the task of *sorting names into alphabetical order*, this Excel spreadsheet has the task of *calculating expected losses*. As well as asking a *how*-question about the device – How does it work? – we might ask a *what*-question: What is the problem it is trying to solve? What is the nature of the task the device faces? A theory at Marr's computational level aims to provide an answer to this question. It aims to identify the *computational problem* that the device faces.¹

What is the computational problem faced by the brain? Conventional approaches in computational cognitive science tend to start from the assumption that the brain faces many distinct computational problems. Different aspects of cognition – e.g. perception, motor control, decision making, language learning – require the brain to respond to different types of computational challenge. Each challenge has its own computational nature and is likely to deserve its own Marrian computational-level description. On such a picture, it makes sense for computational cognitive

¹Marr's use of the term 'computational' here is not meant to imply that his other levels of description are not computational. His usage of the term derives from mathematical logic, where a 'computational' theory denotes relationships between tasks that are blind to differences in algorithms or physical implementation (as in the identification of relations of computational equivalence).

science to adopt a *divide et impera* strategy to modelling cognition: it should break up human cognition into multiple constituent computational problems, each of which should be described in turn.

Predictive coding suggests that this *divide et impera* strategy, and the assumption on which it is based, is wrong. At Marr's computational level, a single, unified story should be told about cognition. During cognition, the brain faces a *single* computational problem. Apparent differences between different challenges that the brain confronts in perception, motor control, decision making, language learning, and so on mask an underlying unity that all these problems share. They are all instances of a single overarching task: to *minimise sensory prediction error*.

Sections 2–4 attempt to unpack what is meant by this claim. Sections 5–8 turn to its justification. I outline three main strategies an advocate of predictive coding might draw on to defend it: the *case-based* defence (Section 7), the *free-energy* defence (Section 8), and the *instrumental-value* defence (Section 9).

2 Minimising sensory prediction error

What does it mean to say that the brain faces the problem of minimising sensory prediction error? As we will see, there are a variety of ways of formalising this task in mathematical language. However, an advocate of predictive coding often starts with an *informal* description of the task. Subsequent mathematical descriptions aim to codify this informal description more precisely and open it up to proposals that it is tackled by various numerical algorithms. There is currently in predictive coding some degree of uncertainty about the right way to formalise the task of minimising sensory prediction error in mathematical terms. However, there is broad agreement about the *informal* nature of the problem. We will begin with this informal description.

The task of *minimising sensory prediction error* may be informally characterised as follows. Brains have sensory organs and their sensory organs supply them with a continuous stream of input from the outside world. Brains also have complicated endogenous physical structures and activities that determines how they react to that stream of input. According to predictive coding, the computational task that a brain faces in cognition is to ensure that these endogenously generated responses (the brain's 'inference' over its 'generative model') cancel out or suppress the incoming flux of physical signals conveyed by the sensory organs from the outside world (that it 'predicts' the incoming 'sensory evidence'). The degree to which this happens, or fails to happen, is measured by the *sensory prediction error*. This quantity measures the discrepancy between the contribution of the brain's endogenously generated activities and the incoming physical signals from the world. According to predictive

coding, the problem that the brain faces, in all aspects of cognition, is to minimise the difference between these two elements. If the brain were to succeed at doing this then, at the sensory boundary the two opposing forces – the world’s sensory input (excitatory/stimulating) and the brain’s endogenously generated predictions (inhibitory/suppressing) – would exactly cancel out. The brain’s anticipatory activity would ‘quench’ the incoming excitation from the world. In more colourful and metaphorical language:

... this is the state that the cortex is trying to achieve: perfect prediction of the world, like the oriental Nirvana, as Tai-Sing Lee suggested to me, when nothing surprises you and new stimuli cause the merest ripple in your consciousness. (Mumford, [1992](#), p. 247, n. 5)

Predictive coding, at least in the first instance, is a theory of the subpersonal computational workings of cognition, not a theory of conscious experience, but the basic idea described in the quotation is sound. The computational task the brain faces during cognition is to avoid being perturbed or surprised by incoming sensory inputs (in the Shannon information-theoretic sense of ‘surprise’, i.e. unpredicted). The brain’s goal is to arrange itself and its physical responses to anticipate and cancel its upcoming sensory input. This goal – ‘Nirvana’ in the above quotation – is unlikely to ever be achieved, or achieved in any sustained way, because the sensory input supplied by the world is almost always too rich and complicated for our brains to be guaranteed to predict it with perfect accuracy. Nevertheless, *trying* to predict it is the task the brain faces in cognition.

Predictive coders suggest that the various computational problems that the brain faces during perception, learning, motor control, decision making, and so on are all instances of this single problem of minimising sensory prediction error. Our various cognitive capacities (sensing, planning, and so on), which have traditionally been viewed as individual solutions by our brain to entirely distinct problems (in perception, motor control, and so on), should be viewed as parts of a seamless, unified response by the brain to a single problem. This suggests that we might need to rethink how we describe and individuate our cognitive capacities, and potentially blur the boundaries between them. Predictive coding aims to offer a grand, unified theory of cognition at Marr’s computational level.

To say that minimising sensory prediction error is *one* of the computational challenges faced by the brain is not novel or unusual. It is common for contemporary models to suggest that the brain engages in compression of sensory signals (Sprevak, [forthcoming\[a\]](#), Section 2) and that certain inference and learning tasks (particularly, in perception) can be described as minimising sensory prediction error (ibid., Section 4). What marks out predictive coding as special is that it says that minimising sensory prediction error is the brain’s *exclusive* computational task. It is not

one problem among many faced by the brain, but its only problem. This elevated role of the task of minimising sensory prediction error is the primary feature that differentiates predictive coding from rival paradigms at Marr’s computational level.

3 Formal and informal descriptions

Theories at Marr’s computational level are often precise and characterised in mathematical language. They are usually *formal* and *quantitative*. Typically, a theory at Marr’s computational level will ascribe computation of a mathematical function to the brain as well as an explanation of why computing that function would help the brain solve a problem as informally characterised. For example, in his account of vision Marr ascribed computation of the mathematical function $\nabla^2 G * I$ to the brain. Marr related this computation to the informally characterised task of *edge detection*: finding the location of boundaries between objects in the visual field.² Marr argued that edge detection is an important task that the brain faces in early vision and that it is a precursor to solving other problems such as object recognition, depth perception, or binocular fusion. Marr proposed that the informal task of edge detection could be more precisely described by formalising it as the task of computing of this mathematical function.

In Marr’s formal description, I is a two-dimensional matrix of numerical values. These numerical values quantify the magnitude of light falling on a two-dimensional array of photoreceptors on the retina. G is a Gaussian filter which is convolved ($*$) with the two-dimensional image (I) and the Laplacian, second-derivative operator (∇^2) is applied to the result. Marr argued that if the brain were to compute the zero-crossings of this function for various sizes of Gaussian filter, it could identify areas in the retinal image that correspond to sharp changes in light intensity. These, Marr argued, tend to coincide with the edges of objects in the visual field. Hence, the task of computing the zero-crossings of this mathematical function provides a precise, mathematically codified formalisation of the (informally characterised) problem of edge detection.³

One way in which this relationship is described is that between a ‘what’ and a ‘why’ element of the computational-level theory.⁴ The ‘what’ element of a computation-

²Marr (1982), pp. 68–74. The full story about the informal task is complex, and ‘edges’ should be understood to include not only the boundaries of objects, but also regions of the visual field where there are changes in reflectance, illumination, depth, or surface orientation.

³Marr thought that this task was solved by the retinal ganglion cells: ‘Take the retina. I have argued that from a computational point of view, it signals $\nabla^2 G * I$ (the X channels) and its time derivative $\partial/\partial t(\nabla^2 G * I)$ (the Y channels). From a computational point of view, this is a precise specification of what the retina does.’ (Marr, 1982, p. 337).

⁴See Marr (1982), p. 22.

level theory describes the mathematical function that the device needs to compute. In the above case, this would be $\nabla^2 G * I$. The ‘why’ element links the task of computing that mathematical function to some informally characterised information-processing problem. It draws a connection between the abstract numerical values that feature in the function and physical quantities in the device and the adaptive problems faced by the embodied device. In the case above, it involves explaining why computing $\nabla^2 G * I$ would help solve the problem of detecting edges in the visual field. Marr’s ‘what’ element provides a formal, mathematical characterisation of the task; the ‘why’ element explains the appropriateness and adequacy of that mathematical description to the task as informally conceived.⁵

There are many possible ways one might attempt to formalise the task of minimising sensory prediction error. Predictive coding has not yet settled on a single canonical or complete way to formalise it. A simple example of a formalisation is given in Sprevak ([forthcoming\[b\]](#)), Section 2.1.⁶ It is worth stressing that contemporary attempts to formalise the task typically aim to formalise a deliberately simplified or stripped-down version of the task as informally characterised. For example, it is normal to only consider systems that only have a few sensory input channels, that only minimise current prediction error, or that make use of simple generative models with linear responses. This is in order to keep the formalisation manageable or in order to illustrate specific features of the intended model.

Nevertheless, some generalisations can be made about predictive coding’s task as formally characterised. All the mathematical formalisations tend to treat the task as an instance of a numerical *optimisation problem*. The optimisation problem is regarded as having two free variables – the *generative model* and the *prediction values*. These should be varied, across different timescales, in order to minimise some objective function – *sensory prediction error*. In the simplest case, the generative model is formalised as a two-dimensional matrix of values. Prediction values are formalised as a vector that, when combined with the generative model by multiplication, produce another vector, the *sensory prediction*. The *sensory input* is a vector with the same dimensionality, each of whose components encode the actual incoming activity in each separate physical sensory channel. The *sensory prediction error* measures how close the prediction is to the sensory input. It is often treated as the (weighted) sum or mean of the squares of the difference between the sensory input vector and the sensory prediction vector. The task the brain faces – as formally characterised – is to select a set of prediction values and a generative model such that its prediction errors over sensory inputs are minimised. Characterising the problem

⁵See Shagrir and Bechtel (2013); Shagrir (2010) for a helpful explanation of the ‘what’ and ‘why’ at Marr’s computational level.

⁶A range of other formalisations can be found in Bogacz (2017); Friston (2003); 1330–1339; Friston (2005), pp. 819–821; Friston (2009), p. 296; Spratling (2017), pp. 92–93.

in this way allows many existing optimisation algorithms, and in particular the vast range of algorithms that involve some form of gradient descent, to be brought to bear as proposals about how the brain attempts to solve its problem.

4 Precision weighting of prediction errors

An important element that has not yet been mentioned is that not all sensory prediction errors matter equally in the task of minimising sensory prediction error. Predictive coding introduces a further variable, *precision weighting*, which describes the relative weight of each sensory prediction error with respect to the others. The brain's task is to minimise *precision-weighted* sensory prediction error. Errors that have a high precision weighting should be prioritised during the optimisation task; errors that have a low precision weighting should be given a lower priority or even partly discounted. Precision weighting thus describes a scaling factor or 'gain' that is applied to each component of the sensory prediction error.

Precision weighting is a critically important part of the task description. It can make certain sensory prediction errors dominate the optimisation process and others small enough to be ignored. It can exercise this control in very fine-grained, nuanced ways. Precision weighting can potentially modify the gain on prediction errors associated with each individual sensory channel independently. Precision weighting is usually treated as a distribution that determines which sensory prediction errors are boosted and which are dampened at any given moment. The shape of that distribution may be complicated and it may change radically and rapidly over time (i.e. within milliseconds). Formally, precision weighting in the simplest case is represented as a two-dimensional matrix that is multiplied by the raw sensory prediction error vector to scale its elements.⁷

Precision weighting plays a number of seemingly distinct roles within predictive coding. First, under a probabilistic interpretation of predictive coding's algorithm, it is assumed to be connected to the brain's estimation of *uncertainty* associated with its sensory predictions. Predictions about which the brain is more confident have a smaller variance, which is equivalent to a greater precision weighting being associated with their corresponding prediction errors (Friston, 2003).⁸ Second, precision weighting is suggested to be connected to the *direction of fit* of sensory predictions. Sensory prediction errors that are assigned a high degree of precision weighting are the ones that the brain is more likely to act on, and hence function as quasi motor commands (see Section 7). Prediction errors that are assigned low precision weighting can be used to simulate or imagine actions of the agent or of

⁷See Sprevak ([forthcoming\[b\]](#)), Section 2.4.

⁸See Sprevak ([forthcoming\[b\]](#)), Section 5.

other agents without generating actual motor responses – they do not meet the threshold of precision weighting to drive motor responses (Clark, 2016; Friston, Mattout and Kilner, 2011, Ch. 5; Pickering and Clark, 2014).⁹ Third, precision weighting is claimed to be connected to the allocation of *attention*. When the cognitive system attends to certain features, the components of the sensory signals associated with those features are the ones for which the corresponding prediction errors have been assigned a higher weight. When the cognitive system shifts the focus of its attention, this entails rebalancing the distribution of precision weightings away from those features (Feldman and Friston, 2010).¹⁰ Finally, and most controversially, precision weighting is sometimes by fans of predictive coding as a kind of ‘fudge factor’ to accommodate data that do not straightforwardly fit into the prediction-error-minimisation task description. If the brain fails to minimise a sensory prediction error, then an advocate of predictive coding might interpret that failure, not as evidence against predictive coding, but as evidence that the brain assigns a low precision weighting to that particular error. If one were to assume an appropriate distribution of precision weightings at each moment in time, almost any observation can be accommodated under predictive coding’s task description.¹¹ Some constraints are needed on how theorists assign precision weightings to the brain. A number of constraints do arise from assumptions made at the algorithmic and implementation levels (Sprevak, forthcoming[b], Section 2.4; Sprevak, forthcoming[c], Section 5), however, finding a sufficient number of empirically motivated constraints on the assignment of precision weightings remains an open problem for predictive coding.¹²

The distribution of precision weighting intuitively captures ‘what matters’ to the brain when it is minimising sensory prediction error. No version of predictive coding can afford to omit this element: it would simply be implausible to say that all sensory prediction errors matter equally to the brain during cognition. However, the introduction of precision weighting into predictive coding’s task description raises a number of puzzles. It plays many roles within predictive coding’s model and it is not obvious all how these various roles cohere. It is also not clear where independent constraints lie on the assignment of precision weightings given its tremendous power to reshape the computational task facing the brain.

⁹See Sprevak (forthcoming[b]), Section 6.1.

¹⁰See Sprevak (forthcoming[c]), Section 5.

¹¹See Clark (2013a) for examples of how precision weighting can explain a range of otherwise puzzling cases (e.g. habit-based action and behaviour during model-free learning). See Miller and Clark (2018), p. 2568 for their response to the objection that precision weighting is a ‘magic modulator’ that allows predictive coding to accommodate every possible goal.

¹²For further discussion of this problem, see Sprevak (forthcoming[c]), Section 8.

5 Long-term prediction error and the dark-room objection

A second important element of the task description that has not yet been mentioned is that, at Marr’s computational level, the objective should be understood as minimising *long-term* sensory prediction error. That goal may be glossed in various ways by advocates of predictive coding with expressions such as ‘global’ prediction error (Lupyan, 2015), ‘upcoming’ prediction error (Muckli, 2010, p. 137), ‘long-term average’ of prediction error (Hohwy, 2013, p. 90, 175, 176), or ‘long-term average surprise’ (Schwartenbeck et al., 2013).

The precise nature of this long-term objective is not entirely clear. Plausibly, it is to minimise the average of individual (precision-weighted) sensory prediction errors over time. However, what type of average, and how far into the future that time should extend, is not clear. It is unknown whether, and to what degree, future prediction errors should be discounted. It is unknown whether the objective should be to reduce prediction errors relative to the system’s own expectations (its subjective probability) of making future sensory prediction errors, or relative to the objective chances (objective probability) of it making such errors. It is unknown whether the relevant extended period is of the order of hours, days, years, the entire future lifespan of the organism, or stretches even further to include the lifespan of its possible descendants and evolutionary successors. It is unknown how the long-term average (which weights prediction errors over time) interacts with precision weighting (which weights the current error signals) – i.e. whether precision weighting should be understood as having a prospective component. These open questions suggest that alternative formulations of predictive coding could be developed at the computational level.

Nevertheless, acceptance that the brain aims to minimise a long-term measure of prediction error plays an important role in clarifying and lending plausibility to predictive coding’s task description, even if the exact nature of that long-term objective is not clear. It allows one to understand how predictive coding can respond to the infamous ‘dark room’ objection. It also suggests that predictive coding is compatible with inferences and behaviour that tend to drive up short-term sensory prediction error, such as curiosity, exploration, and novelty seeking.

The dark-room problem is a long-standing objection to predictive coding.¹³ The problem is to explain why, if predictive coding’s computational-level description is true, cognitive agents like ourselves do not simply seek out the most predictable possible environment, such as a dark room, and remain inside for as long as possible. If the goal of cognition is only to minimise sensory prediction error, why not maximise the chances of this happening by choosing to stay in a maximally

¹³See Clark (2013b), p. 193 for a statement of the problem.

predictable environment?

Friston, Thornton and Clark (2012) offered an initial reply to the dark-room problem.¹⁴ Their response focused on the idea that our generative model and prediction values, as physically implemented in our neural hardware, are not infinitely malleable. There are limits to the kinds of predictions we can generate and to how much our generative model and prediction values can be revised – these constraints, which are assumed to be immune to change by learning or inference, are called ‘hyperpriors’. The kinds of sensory data that a hypothetical cognitive system might receive inside a dark room may be predictable in some abstract sense but, due to the peculiar nature of our hyperpriors, that data might be difficult for creatures *like us* to predict. A different type of organism, one with different hard-wired biases (maybe a cave-dwelling creature), might have no trouble in generating accurate predictions inside a dark room. However, humans are strongly biased to predict sensory changes, and so we are unlikely to minimise our sensory prediction errors inside a dark room.

This response highlights an important and as yet unmentioned point about predictive coding’s computational-level claim: the task facing the brain is a *constrained* optimisation problem. The goal of the brain is to minimise its sensory prediction errors by varying its generative model and its prediction values *given* the constraints imposed by our physical hardware about how far and how rapidly that generative model and those prediction values can vary. The brain aims to minimise its sensory prediction errors relative to a variety of physical constraints. Predictive coding, at the computational level, tends to leave details about the nature of those constraints largely unspecified.¹⁵

Even if this reply is correct, one might worry that it does not fully address the concerns that motivated the dark-room objection. For example, it does not explain why, *even relative to a constrained model*, cognitive agents like ourselves still seek out novelty and surprise. Even if we *can* accurately predict a situation, we sometimes choose to shun it for a more surprising alternative. In other words, cognitive agents like ourselves sometimes *prefer* novelty and surprise to predictability. How is that behaviour consistent with what predictive coding says at the computational level?¹⁶

An alternative reply, which fares better at addressing this kind of objection, is to

¹⁴See also Hohwy (2013), pp. 87, 185; Clark (2016), pp. 265–268;

¹⁵We will see that some information about the constraints flows from what predictive coding says at the algorithmic level and implementation level (Sprevak, forthcoming[b], Section 2.5; Sprevak, forthcoming[c], Section 7). However, as will become clear, what predictive coding says at those levels is by no means either a complete or a settled account of the relevant constraints faced by the brain in inference or learning.

¹⁶See Clark (2016), pp. 265–266

emphasise the long-term nature of the brain's prediction-error-minimisation task. The world in which we live contains both environments that are easy to predict and environments that are hard to predict. Successfully predicting our sensory inputs only where we can already do so may not, over the long term, be a good solution to the brain's problem. An agent who sequesters itself inside a dark room or a similar predictable environment leaves itself a hostage to fortune. Unpredictable elements may impose themselves on the agent in ways that it has not taken the trouble to learn – light might enter the room, a stranger might enter, food supplies might run out. To guard against possible unpleasant future surprises and the associated rise in sensory prediction error, it may be better – in the terms of meeting the long-term goal of minimising sensory prediction error – to leave the dark room now and engage in some exploration to learn a more comprehensive model of the world – that would allow the agent to predict a broader range of scenarios. Exploring now might raise sensory prediction errors in the short term, but it is a hedge against future surprises that an agent who leads an entirely sheltered life would not be able to predict. There is obviously a balance to strike between the cost of acquiring this information (in terms of an expected rise in short-term sensory prediction error), and its potential future pay-off (in terms of an expected reduction in long-term sensory prediction error). But that there is a trade-off between the adaptive value of exploration and exploitation is to be expected on any model of cognition. The important point is that what predictive coding says at the computational level is compatible with cognitive agents sometimes preferring unpredictable environments to predictable ones. Curiosity, exploration, and novelty seeking are consistent with the brain minimising a long-term measure of sensory prediction error, even if they entail a short-term increase in that error along the way (Schwartenbeck et al., 2013).

6 Evidence for predictive coding

Justification for predictive coding's computational-level claim often rests on one of three strategies. I call these strategies the *case-based* defence, the *free-energy* defence, and the *instrumental-value* defence. The case-based defence considers a range of cognitive tasks and aims to show that all of these tasks can and should be described as minimising sensory prediction error. The free-energy defence shortcuts consideration of individual tasks and attempts to establish predictive coding's computational-level claim in one fell swoop by appeal to Karl Friston's 'free-energy' principle. The instrumental-value defence focuses on the utility of predictive coding to computational cognitive science and argues that it provides a desirable set of heuristics to make sense of, and discern patterns within, the mass of human behavioural and neural responses.

7 The case-based defence

The case-based defence is an abductive argument. It attempts to show that a number of tasks facing the brain – for example, during perception, decision-making, planning, motor control – can and should be thought of as instances of the single task of minimising sensory prediction error. Some of those tasks may already have computational-level descriptions associated with them based on rival or more traditional computational research programmes. The job of predictive coding is to show that these should be reconceptualised as instances of minimising sensory prediction error. Behavioural and neural responses that might previously have been described as attempts by the brain to compute some domain-specific mathematical function should be redescribed in the manner predictive coding suggests.

Any case-based argument for predictive coding faces an obvious epistemic hurdle. Predictive coding makes a universal claim – *every* problem the brain encounters in cognition is to minimise sensory prediction error. Showing that this claim holds in some cases (e.g. for early vision) will not entail that it holds in other, perhaps as yet unconsidered cases (e.g. for language learning). No consideration of individual cases entails the conclusion that in *every* case the problem the brain faces can and should be described as minimisation of sensory prediction error. Nevertheless, science is rife with universal generalisations made on the back of observations about specific cases. The non-demonstrative nature of such arguments is not an in principle objection to using them. However, there are clearly more and less effective ways of making such an argument work.

One plausible strategy for making the universal generalisation credible is to focus on a *diverse* range of cases – what one might hope is a *representative* sample of what the brain is up to in cognition. Early work on predictive coding focused on sensory compression in the early visual system (Atick, 1992; Rao and Ballard, 1999; Srinivasan, Laughlin and Dubs, 1982). Ideally, predictive coding should seek support for its claim by showing that other kinds of behavioural and neural response fall under predictive coding’s task description. If it can be shown that many, diverse behavioural and neural phenomena that have no obvious connection to each other, or to the early visual system, can and should fall under predictive coding’s task description, then that would lend credence to the abductive generalisation that not just in some cases, but in every case, the problem the brain faces is sensory prediction error minimisation. Example of ‘non-obvious’ applications of predictive coding include music perception (Koelsch, Vuust and Friston, 2019); formation of emotions and judgements about bodily ownership (Seth, 2013); binocular rivalry (Hohwy, Roepstorff and Friston, 2008); formation of judgements about the nature of the self (Hohwy and Michael, 2017); and the perceptual, doxastic, and motor characteristics of schizophrenia and autism (Corlett and Fletcher, 2014; Fletcher

and Frith, 2009; Friston, Stephan et al., 2014; Pellicano and Burr, 2012).

It is worth noting that, for each individual case, a case-based argument requires one to meet two separate challenges. The first challenge is to show that the case in question *can* be described as an instance of sensory-prediction-error minimisation. The second is to show that it *should* be described this way. The first challenge requires one to show that predictive coding's computational-level description is *consistent* with the behavioural or neural data associated with that case. The second challenge is to show that one should *prefer* predictive coding's computational-level description of that data to rival or traditional accounts. There should be some benefit to adopting predictive coding's computational-level treatment of that instance of cognition – e.g. in terms of increased predictive accuracy, increased explanatory power, or some other theoretical virtue.

Predictive coding's flagship example of a 'non-obvious' case is *motor control*.¹⁷ Traditional approaches tend to categorise perception and motor control as separate problems. In perception, the task facing the brain is to use its sensory data and background knowledge to build an accurate (or adequate) model of the world. In motor control, the task facing the brain is to use that perceptual model, along with some set of goals or intentions, to output a sequence of motor commands that will direct muscle actuators towards accomplishing those goals or intentions. Of course, motor control might partly involve solving a perceptual problem. Motor problems often require an agent to first build an adequate perceptual model to guide motor planning. Rapid and complex motor control might require regulation by sensory predictions from a forward model (Franklin and Wolpert, 2011). However, even if the problems of motor control and perception have some overlap, they remain separate problems: the objective of perception is to create an accurate model of the world; the objective of motor control is to use that model to generate motor commands to fulfil goals.

According to predictive coding, perception and motor control should be conceptualised as exactly the same problem, namely, to minimise sensory prediction error. In perception, the brain minimises sensory prediction error by varying its endogenous generative model and prediction values to yield predictions that minimise error over its incoming sensory stream. In motor control, the brain minimises sensory prediction error by varying its bodily position and the external world (via muscle actuators) to modify its incoming sensory stream to make the endogenously generated sensory predictions true. In both cases, the objective is the same – to minimise sensory prediction error. The difference lies in the method the cognitive system uses to try to achieve it. Advocates of predictive coding call the first method

¹⁷See Friston (2010), pp. 133–134; Friston, Daunizeau et al. (2010); Clark (2016), Section 4.5; Hohwy (2013), Ch. 4.

‘passive’ inference and the second ‘active’ inference. Passive and active inference (perception and motor control) are complementary strategies employed by the brain to address what is fundamentally the same problem. According to predictive coding, the task of reaching for a glass of water should be reconceptualised as the brain making the prediction that the hand is already holding the glass of water (along with all its sensory consequences), and then solving the problem – minimising its sensory prediction error – by varying its limbs and the glass to make the false sensory prediction true.¹⁸

What is mooted here is that perceptual tasks and motor control tasks *can* both be described as instances of sensory prediction error minimisation. Even if this is true however, it remains a further question whether they *should* be described this way. The justification given for this second claim is often not obvious. The benefits of predictive coding’s task description are not straightforward to calculate and need to be measured relative to a broad range of epistemic standards, interests, and goals in computational cognitive science. Different research groups may take different views about the value of the benefits on offer.¹⁹ As we will see, the benefits are also often conditional on accepting other elements of predictive coding’s research programme (e.g. the universal scope of its claim, or elements of its proposals at the algorithmic and implementation levels).

To illustrate how these questions about preferability might be addressed, we will switch to a simpler case: the early visual system. Two main strategies have been used to justify predictive coding’s computational-level description in this context: (i) its predictive and explanatory benefits over traditional computational approaches; (ii) the broader theoretical virtues offered by the view (e.g. simplicity, elegance, and unifying power).

The first set of considerations surround predictive coding’s ability to predict and explain individual behavioural or neural responses that are puzzling or anomalous on other views. Traditional computational-level characterisations of the early sensory system suggest that its task is to act as a Gabor filter bank on retinal images to extract ecologically salient stimulus features such as orientation, spatial frequency, colour, direction of motion, and disparity (Carandini, Demb et al., 2005). The computational problem faced by neurons in the early visual system is to convolve a matrix of retinal data with a range of Gabor filters to, e.g., pick out lines in the visual

¹⁸Predictive coding also provides an *algorithmic-level* proposal about how motor tasks are solved. As we will see, the suggested algorithms for perception and motor control have a great deal in common (see Sprevak, [forthcoming\[b\]](#), Section 6.1).

¹⁹For the benefits of predictive coding’s task description of motor control see Friston (2011), Friston, Daunizeau et al. (2010); Wiese (2017), Pickering and Clark (2014). For benefits of alternative approaches, see Kording (2007); Shadmehr and Krakauer (2008).

field of various orientation and spatial frequency. However, many physical responses exhibited by the early visual system do not fit this computational-level description (Olshausen and Field, 2005). One such ‘non-classical’ effect is *end-stopping*: neurons in V1 give a strong response to a line at a particular orientation in the visual field, but this response is reduced or eliminated if the line extends outside that neuron’s receptive field. End-stopping is inconsistent with a simple Gabor-filter description of their computational role. A classical Gabor filter should continue to fire regardless of whether a line extends outside its receptive field. End-stopping counts as an unexplained anomaly under traditional computational-level description of the early visual system.

Predictive coding suggests that the function of the early visual system is to contribute to minimising the cognitive agent’s sensory prediction error. Under predictive coding’s task description, the behaviour of neurons within V1 may be reinterpreted as signalling the difference between the current sensory input and the brain’s sensory prediction (based on its statistically-informed expectations regarding visual input). In our environment, the statistical norm is for lines in the visual field to extend beyond the tiny regions covered by individual receptive fields. Lines that violate this expectation are unusual and, everything being equal, should be expected to cause sensory prediction errors. The behaviour of V1 cells when end-stopping may be interpreted as signalling such sensory prediction errors (Kok and de Lange, 2015; Rao and Ballard, 1999, p. 232). End-stopping is an anomaly on traditional computational-level descriptions, but it can potentially be predicted and modelled on predictive coding’s computational-level description.²⁰

A second set of motivations for preferring predictive coding’s task description surround predictive coding’s general theoretical virtues such as its simplicity, scope, and unifying power with respect to other computational-level approaches. Arguably, even if predictive coding does no better than alternative approaches in terms of modelling anomalous behavioural/neural effects, those general virtues might still lead one to favour the view. As observed in Section 1, traditional computational-level approaches to cognition tend to categorise the brain as facing multiple, largely unrelated computational problems. This may lead one to assume that the brain is an inherently multifunctional device, rather than a device tuned to solve just one problem.²¹ A description of human cognition at Marr’s computational level may

²⁰For other examples of non-classical effects in the early visual system that appear to be predicted and modelled by predictive coding, see Jehee and Ballard (2009); Kok, Jehee and de Lange (2012); Hosoya, Baccus and Meister (2005); Rao and Sejnowski (2002); Muckli (2010); Kok and de Lange (2015); Spratling (2010); Alink et al. (2010); Murray et al. (2002). For alternative computational-level accounts of these non-classical phenomena (e.g. in terms of divisive normalisation), see Aitchison and Lengyel (2017), p. 224; Carandini and Heeger (2012); Schwarz and Simoncelli (2001).

²¹For example, see Allen (2017); Bayne et al. (2019).

thus be expected to consist in a patchwork of disjoint theories covering each task the brain faces. Each task facing the brain – perception, motor control, decision making, language learning – may merit its own computation-level account. Stepping back from this patchwork, there need be no overarching pattern or unity to cognition and plenty of gaps in our existing accounts of it. Predictive coding, in contrast, provides a unified, complete, and relatively simple description of the computational problem the brain faces in all aspects of cognition. That by itself appears to be a mark in its favour.

It is the first time that we have had a theory of this strength, breadth and depth in cognitive neuroscience ... I take that property as a sure sign that this is a very important theory ... Most other models, including mine, are just models of one small aspect of the brain, very limited in their scope. This one falls much closer to a grand theory. (Stanislas Dehaene quoted in Huang, 2008)

A unified computational-level theory promises to reveal something profound about the fundamental nature of cognition. It tells us that cognition is not a motley, a jumble of distinct phenomena, but a response to a single computational problem. Predictive coding identifies what the various, seemingly distinct and unrelated departments of human cognition – e.g. perception, motor control, decision making, language learning – have in common. It purports to explain why they each count as instances of cognition. It provides us with information to judge whether new and perhaps unexpected or previously unconsidered instances of cognition are genuinely cognitive.²² Moreover, predictive coding suggests that cognition is in essence a unified and simple functional kind. If a theory uncovers fundamental principles like this, that unifies and simplifies an otherwise disordered domain, then, everything else being equal, that is a reason to favour it. Knowledge about the essence of things and general patterns into which they enter is surely what science aspires to.

8 The free-energy defence

Any case-based defence of predictive coding is likely to be a long project and fraught with difficulties. It requires engaging with the details of many different individual tasks and showing that their distinctive effects – of which there may be many – are captured or recaptured on predictive coding's task description. The defence also has no obvious stopping point at which victory could be declared. A defender of predictive coding faces a potentially endless sequence of battles: there will always

²²For predictive coding as a potential 'mark of cognitive', see Clark (2017); Kirchhoff and Kiverstein (2019); Ramstead et al. (2021).

be more tasks, more behavioural and neural effects to consider, in order to argue for the merits of describing the problem the brain faces in terms of predictive coding. It is not obvious when enough cases – or a diverse enough array of cases – will have been considered to warrant the conclusion that not just *some* tasks, but *every* task faced by the brain is sensory prediction error minimisation.

The free-energy defence aims to shortcut this. It tries to establish predictive coding's computational-level claim in a single step by appeal to general properties shared by all cognitive (or living) systems. Friston (2010) presents a defence of predictive coding along these lines based on his 'free energy' formulation of predictive coding. Friston proposes that the task faced by the brain is that of *minimising free energy*. Minimising free energy can be shown, if appropriate further assumptions are made, to be equivalent to the task of minimising sensory prediction error.

Free energy is a mathematical quantity that appears in classical thermodynamics, statistical mechanics, and information theory. Friston's central claim is that there is a relationship between two distinct applications of the free-energy formalisation: *variational* free energy and, what I will call, *homoeostatic* free energy.²³ Variational free energy is an information-theoretic quantity predicated of agents who engage in probabilistic inference. If a probabilistic reasoner minimises their variational free energy, then this can be shown to be equivalent to them approximating Bayesian inference (see Sprevak, forthcoming[d], Section 1). Granted a number of further assumptions, minimising variational free energy can also be shown to entail minimising sensory prediction error (see Sprevak, forthcoming[d], Section 2). 'Homoeostatic' free energy is a distinct quantity which applies the same abstract free-energy formalism to an entirely distinct set of properties. Unlike variational free energy, it is not (or at least, not directly) associated with the subjective probabilities that feature in an agent's probabilistic inferences. Rather, it is associated with the objective probability of the macroscopic physical state the agent is in given its physical environmental conditions. Minimising homoeostatic free energy is associated with the agent's survival within a narrow band of macroscopic physical states. According to Friston, these two kinds of free energy – homoeostatic free energy and variational free energy – are interlinked. Agents who minimise their homoeostatic free energy – i.e. who survive and maintain homeostasis and thereby maintain their macroscopic physical state in response to likely environmental perturbations – also minimise their variational free energy (and hence, given certain assumptions, minimise their sensory prediction error).

Friston is clear that the free-energy quantity he has in mind is not the same as *thermodynamic* free energy. Thermodynamic free energy intuitively measures the 'useful' work that can be obtained from a physical system. It is usually defined in terms of

²³Friston does not use these terms. He refers to both as 'variational' free energy.

that system's ability to exert macroscopic mechanical forces on its surroundings – its energy that is 'free' to perform mechanical work. This is normally formalised as a difference between the physical system's 'internal energy' and its thermodynamic entropy – its energy that is 'useless' for work. Having a reserve of thermodynamic free energy is generally a useful resource for a cognitive or living creature: having thermodynamic free energy is a prerequisite for the creature to be able to move or act in the world. *Minimising* thermodynamic free energy would make little sense as a survival strategy or as a way to maintain physiological functioning. Friston is explicit that his free-energy principle – that all cognitive/living systems aim to minimise their homoeostatic/variational free energy – is not meant to be somehow a consequence of thermodynamics or a principle about thermodynamic free-energy. His free-energy principle is instead justified on 'selectionist' grounds: all cognitive/living creatures strive to minimise their homoeostatic free energy because if they did not do so, they would tend to die off and hence be less likely to reproduce or to be observed by us.²⁴ Friston suggests that the only connection between thermodynamic free energy and his notion of free energy is their shared mathematical form.²⁵

In outline, the logic of the free-energy defence of predictive coding is as follows. Its starting point is the observation that all cognitive (and living) creatures face the problem of surviving and maintaining homeostasis. That task, according to Friston, can be formally redescribed as the task of minimising a free-energy measure (what I have called homoeostatic free energy). Friston claims that minimising this free-energy measure entails that the creature also minimises a second free-energy measure associated with the creature's subjective probabilistic guesses (variational free energy). Minimising variational free energy, given certain further assumptions (detailed in Sprevak, [forthcoming\[d\]](#), Section 2), entails that the creature also minimises its sensory prediction error. Hence, cognitive and living creatures, because they face the problem of survival and maintaining homeostasis, face the problem of minimising sensory prediction error.

There is much to unpack here.

First, the argument relies on a presumed connection between homoeostatic and variational free energy. However, the justification for that connection is not obvious. Homoeostatic free energy pertains to how well the creature maintains its physical state within the narrow band associated with survival and homeostasis in the face of actual and possible perturbations from a changing physical environment. Living creatures change their microscopic physical state all the time. When they do so, they risk undergoing a fatal phase transition in their macroscopic physical state. When

²⁴Friston and Stephan (2007), pp. 419–420, 451; Friston, Kilner and Harrison (2006), p. 85

²⁵See Friston and Stephan (2007), p. 419.

living systems resist this tendency – when they survive and maintain homeostasis – they minimise their homoeostatic free energy. Minimising homoeostatic free energy involves the creature trying to arrange its macroscopic states so as to avoid being overly changed by likely environmental physical transitions. A physical system that minimises its homoeostatic free energy strives to maintain its macroscopic physical state in equipoise with likely environmental changes (Friston, 2013; Friston, Kilner and Harrison, 2006; Friston and Stephan, 2007). In contrast, variational free energy is an information-theoretic quantity predicated of an agent’s subjective probability distributions. It measures how far the agent’s probabilistic guesses depart from the optimal guesses of a perfect Bayesian observer armed with the same evidence.²⁶ According to Friston’s formulation, the brain’s task is to minimise this variational free-energy quantity and so approximate an ideal Bayesian reasoner in inference. Minimising variational free energy makes the sensory data stream unsurprising (in the information-theoretic sense), and thereby tends to minimise the agent’s sensory prediction error (modulo certain assumptions outlined in Sprevak, forthcoming[d], Section 2).

Homoeostatic free energy and variational free energy have certain features in common: they are both information-theoretic quantities and they both attach to probability distributions. However, they are not the same quantity. Homoeostatic free energy is measured over the *objective* probability distributions of macroscopic physical states that could occur; variational free energy is measured over the *subjective* probability distributions entertained by an agent about what could occur. Variational free energy attaches to subjective probability distributions; homoeostatic free energy attaches to chances of various possible (fatal) physical states of the agent occurring in response to environmental changes. The respective probability distributions might in principle be defined over distinct sets of events, their distributions might take different shapes, and they each involve materially different types of probability (subjective and objective). There may be correlations between the two free-energy measures, but it is not obvious that minimising free energy for one probability distribution entails minimising free energy for the other.²⁷

To see this point more clearly, consider the relationship already mentioned between variational free energy and Bayesian inference. An agent who minimises its variational free energy approximates an ideal Bayesian reasoner. In many circumstances a Bayesian agent is well placed, and in some circumstances it will be better placed than a non-Bayesian agent, to survive and maintain homeostasis. But the precise nature of the connection between *being Bayesian* in one’s reasoning and *maximising*

²⁶See Sprevak (forthcoming[d]), Section 1 for the connection between variational free energy and Bayesian inference.

²⁷Sprevak (2020), pp. 602–604.

one's chances of physical survival and homeostasis is far from obvious. A non-Bayesian agent might live in a 'irrational friendly' environment that maintains its homeostasis and physical integrity, even if it does not update its subjective probability distributions which represent that environment according to Bayesian norms. Conversely, an ideal Bayesian reasoner might live in a 'rationality hostile' physical environment that changes so rapidly and dramatically that it fails to survive or maintain homeostasis, even if it updates its subjective probability distributions quickly and represents the environment accurately according to the Bayesian norms. Bayesian reasoning is not unrelated to survival, but it is not obvious in what sense it would guarantee it. In information-theoretic terms, the exact nature of the relationship between Friston's two types of free energy – homeostatic and variational – is unclear and the subject of ongoing analysis.²⁸

At least two other aspects of the free-energy defence of predictive coding invite further scrutiny.

First, the aim of the predictive coding research programme is to defend the claim that *every* task that the brain faces can and should be described as minimisation of sensory prediction error. Survival/homeostasis is clearly one important task faced by a brain. If the internal logic of the free-energy defence is correct, then because the brain faces that task it also faces the task of minimising sensory prediction error. But it is not obvious that survival/homeostasis is the *only* problem faced by a brain. Plausibly, the human brain faces other challenges that may be unrelated, or even in tension with, the human agent's survival or homeostasis – e.g. problems of mate selection, fulfilment of social roles, or arbitrary challenges set in the classroom or wider social environments. It is not clear how the free-energy defence is intended to handle these cases. The defence appeals to the connection between survival/homeostasis and minimising sensory prediction error, but it is largely silent about how problems that do not (or do not obviously) contribute to survival/homeostasis are meant to be related to sensory predictive error. Consequently, even if one were to assume that the internal logic of the free-energy defence is correct, it is unclear how it would establish predictive coding as a universal claim.

Second, recall that the case-based defence required one to show, not only that every computational problem faced by the brain in cognition *can* be redescribed as sensory prediction error minimisation, but also that it *should* be described that way. The free-energy defence only appears to speak to the first issue. It attempts to establish a connection between the task of survival/homeostasis and the task of minimising sensory prediction error. However, even if such a connection exists, it would say nothing about the merits of one task description over the other at Marr's computational level. In order to address that, one would need to go beyond the

²⁸See Bruineberg, Kiverstein and Rietveld (2018); Colombo and Wright (2018); Sprevak (2020).

relationships between tasks as conjectured by the free-energy defence and consider the *value* of predictive coding's proposed redescription with respect to the wider standards, interests, and goals in cognitive neuroscience. *Why* should we describe the task facing the brain as sensory prediction error minimisation, even if, as the free-energy defence suggests, we can? That argument remains to be made, and doing so is likely to depend, at least partly, on an examination of the benefits offered by predictive coding's proposed description at the level of the treatment of specific cases of interest to cognitive neuroscience. This suggests that the free-energy defence may not be able to entirely shortcut the exigencies of the case-by-case defence.

9 The instrumental-value defence

The instrumental-value defence has an entirely distinct character from the previous two. This third strategy for defending predictive coding helps to explain an otherwise puzzling phenomenon: the widespread adoption of its computational-level claim in cognitive neuroscience despite what we have seen as the view's current relatively slender epistemic support. According to the instrumental-value defence, predictive coding should be interpreted, not as a passive claim that awaits confirmation, but as a *discovery heuristic* – an assumption that researchers might adopt in order to help better organise data, guide experimental design and interpretation, and formulate further, more specific hypotheses for testing. Predictive coding's computational-level description provides a novel way to describe and systematise behavioural and neural responses. It constrains the way one might group behavioural and brain responses into psychologically relevant, computationally-defined capacities. Furthermore, if one understands predictive coding as a package that includes proposals at Marr's algorithmic and implementation levels, then it provides a rich set of heuristics to guide and inspire claims about the formal methods and neural mechanisms that underlie those computational capacities. The focus in the previous two sections was on whether predictive coding gets the computational-level description of the brain *right* or *wrong* (or whether it fares better than alternative proposals). But one might equally well ask the prior question of how one can come up with a computational-level description of the brain *at all*. Scientific work here can potentially benefit from what predictive coding says, even if uncertainty remains about the view's ultimate epistemic standing.

Individuating behavioural and neural responses into the exercise of a set of well-defined neural computational capacities is hard. Cognitive neuroscientists do not have an agreed methodology to guide them through this process. Formulating a computational-level description of the brain usually requires adopting some broad theoretical orientation about the overall purpose of the brain's physical activity. It is not obvious where an empirically minded researcher should look to for this.

Traditionally, folk psychology has provided one possible source of inspiration. One might, for example, start by assuming that the brain is trying to use ‘belief’-like and ‘desire’-like states to produce outcomes that satisfy what it represents as ‘desired’. Bringing this to bear on empirical data might motivate a researcher to formulate more specific hypotheses about different kinds of belief-like and desire-like states inside the brain, the relationships between them, the processes that transform them, and how sensory and behavioural responses update those beliefs and fulfil those desires.²⁹

An alternative source of inspiration might lead a researcher towards a different set of specific, testable hypotheses about the computational tasks the brain faces and its underlying computational capacities, states, and mechanisms. Machery ([forthcoming](#)) describes how one feature of evolutionary psychology is that, irrespective of its other epistemic properties, it provides a potentially valuable set of discovery heuristics. One of those heuristics (the ‘forward-looking’ heuristic) speaks directly to the problem of coming up with hypotheses at Marr’s computational level. It suggests that our computational capacities should be identified by looking at the information-processing problems encountered by our ancestors that regularly bore on their fitness.³⁰ The computational capacities that our brains have today should be inferred from the problems faced by our evolutionary ancestors (Cosmides and Tooby, 1989). Hypotheses about our computational capacities arrived at in this fashion of course need to be empirically confirmed. But even in advance of securing epistemic support, it may make sense to accept a framework like evolutionary psychology (or folk psychology) as a discovery heuristic, in order to make the problem of task description tractable at all.

Predictive coding potentially plays a similar role for cognitive neuroscience. It suggests that neural and behavioural responses should be organised around the central idea that those responses are all attempts by the brain to minimise (long-term, precision-weighted) sensory prediction error. Even if the evidential basis for that idea is relatively slim, it may function as a useful heuristic to guide design of experiments, measurement, and generate more specific, testable proposals about physical responses.

For example, Fletcher and Frith (2009), inspired by predictive coding’s computational-level claim, hypothesise that a range of positive symptoms of schizophrenia – including hallucinations, delusions, abnormal saliences in perception, disturbances in low-level motor functioning – should be categorised as instances of a single, unified dysfunction in the computational function to minimise (precision-weighted) sensory prediction error. They go on to propose that

²⁹See Machery ([forthcoming](#)), Section 1.1.

³⁰*ibid.*

this dysfunction is unwritten by a single, unified computational mechanism and physical basis, again prompted by predictive coding's claims at those levels.³¹ Such work suggests novel experimental designs that might attempt to dissociate these factors, probe how they might be quantitatively affected by manipulating sensory prediction errors, and explore analogues of schizophrenia in healthy subjects by looking at regimes that have similar effects on sensory prediction errors.³² Corlett and Fletcher (2014) describe how predictive coding could function as a discovery heuristic for clinicians to find and trial new therapeutic interventions for patients (including pharmacological treatments). The idea that the brain aims to minimise sensory prediction error might function as the starting point for any number of concrete theoretical, experimental, and therapeutic developments.

In contrast to both the case-based defence and the free-energy defence, the focus here is not primarily on truth, but on predictive coding's *utility*. The relevant kind of utility should be understood as broader than merely a concern with achieving a narrowly instrumental outcome. Cognitive neuroscientists need to make assumptions regarding the overall purpose of neural functioning in order to make any sense of activity in the brain and behaviour. Those assumptions need to come from somewhere. It is reasonable that any candidate source for those ideas should be understood to be uncertain and exploratory; predictive coding provides one among many possible approaches (distinct from folk psychology or evolutionary psychology). Its sheer novelty – predictive coding's ability to depart from traditional categorisations of behaviour and neural response – is undoubtedly an attraction. It allows us to see familiar behavioural and neural responses in a new light and group them together in different ways from previous research programmes. The central idea that generated these hypotheses may ultimately prove to be mistaken, but that possibility should not disbar it from being used to guide current thinking or practice.

Using predictive coding in this way – as a heuristic to guide discovery rather than a claim that passively awaits confirmation – does not somehow magically confer justification on the view. Merely believing something does not make it true. Justification for predictive coding only accrues if it can predict and explain empirical results better than alternative theoretical approaches.³³ The instrumental-value defence does not obviate the need to gather empirical evidence to confirm predictive coding. However, it does explain why someone might be rational to accept what predictive coding says now, even in advance of such evidence being obtained. It explains why predictive coding might be adopted in cognitive neuroscience as a working

³¹ibid., pp. 53–55; Corlett, Frith and Fletcher (2009).

³²For example, see Fletcher and Frith (2009), p. 55–56.

³³See Machery (forthcoming), Section 3.2 for a similar point regarding evolutionary psychology.

hypothesis despite its truth remaining in question.

10 Conclusion

This paper has focused on what predictive coding says at Marr’s computational level. In its boldest form, predictive coding proposes that the *only* computational problem that the brain faces in cognition is to minimise its long-term, prediction-weighted sensory prediction error. This paper has reviewed three strategies to defend this claim (Sections 7, 8, 9). These three defences should not be viewed as mutually exclusive, but as potentially complementary methods for justifying predictive coding.

It is natural to wonder what would happen if one were to trim predictive coding’s ambitions.³⁴ Perhaps it describes *some* of the problems that the brain faces, but not all. One might imagine a variety of ways in which its computational-level claim might be reigned in. At one limit would be the relatively uncontroversial claim that *one* thing the brain does, in early vision, is to minimise sensory prediction error to compress sensory signals. At the other end would be the unqualified claim that minimising sensory prediction error is the *only* thing that the brain does. An advocate of predictive coding might wish to adopt a view that falls between these two extremes. However, it is worth noting that the extent to which caveats and qualifications are introduced, the distinctive scope and unifying power of the predictive coding framework is compromised. The predictive coding research programme, if it is to fulfil its original promise, should aim to deliver as broad and comprehensive a theory of cognition as possible.

Bibliography

- Aitchison, L. and M. Lengyel (2017). “With or without you: Predictive coding and Bayesian inference in the brain”. In: *Current Opinion in Neurobiology* 46, pp. 219–227.
- Alink, A., C. M. Schwiedrzik, A. Kohler, W. Singer and L. Muckli (2010). “Stimulus predictability reduces responses in primary visual cortex”. In: *Journal of Neuroscience* 30, pp. 2960–2966.
- Allen, C. (2017). “On (not) defining cognition”. In: *Synthese* 194, pp. 4233–4249.
- Atick, J. J. (1992). “Could information theory provide an ecological theory of sensory processing?” In: *Network: Computation in Neural Systems* 3, pp. 213–251.

³⁴See Clark (2013b), pp. 200–201.

- Bayne, T., D. Brainard, R. W. Byrne, L. Chittka, N. Clayton, C. Heyes, J. Mather, B. Ölveczky, M. Shadlen, T. Suddendorf and B. Webb (2019). “What is cognition?” In: *Current Biology* 29, R603–R622.
- Bogacz, R. (2017). “A tutorial on the free-energy framework for modelling perception and learning”. In: *Journal of Mathematical Psychology* 76.198–211.
- Bruineberg, J., J. Kiverstein and E. Rietveld (2018). “The anticipating brain is not a scientist: The free-energy principle from an ecological-enactive perspective”. In: *Synthese* 195, pp. 2417–2444.
- Carandini, M., J. B. Demb, V. Mante, D. J. Tolhurst, Y. Dan, B. A. Olshausen, J. L. Gallant and N. C. Rust (2005). “Do we know what the early visual system does?” In: *Journal of Neuroscience* 25, pp. 10577–10597.
- Carandini, M. and D. J. Heeger (2012). “Normalization as a canonical neural computation”. In: *Nature Reviews Neuroscience* 13, pp. 51–62.
- Clark, A. (2013a). “The many faces of precision (Replies to commentaries on “Whatever next? Neural prediction, situated agents, and the future of cognitive science”)”. In: *Frontiers in Psychology* 4, p. 270.
- (2013b). “Whatever next? Predictive brains, situated agents, and the future of cognitive science”. In: *Behavioral and Brain Sciences* 36, pp. 181–253.
- (2016). *Surfing Uncertainty: Prediction, Action, and the Embodied Mind*. Oxford: Oxford University Press.
- (2017). “How to knit your own Markov blanket”. In: *Philosophy and Predictive Processing*. Ed. by T. Metzinger and W. Wiese. Frankfurt am Main: MIND Group. DOI: [10.15502/9783958573031](https://doi.org/10.15502/9783958573031).
- Colombo, M. and C. Wright (2018). “First principles in the life sciences: The free-energy principle, organicism, and mechanism”. In: *Synthese*. DOI: [10.1007/s11229-018-01932-w](https://doi.org/10.1007/s11229-018-01932-w).
- Corlett, P. R. and P. C. Fletcher (2014). “Computational psychiatry: a Rosetta Stone linking the brain to mental illness”. In: *The Lancet Psychiatry* 1, pp. 399–402.
- Corlett, P. R., C. D. Frith and P. C. Fletcher (2009). “From drugs to deprivation: a Bayesian framework for understanding models of psychosis”. In: *Psychopharmacology* 206, pp. 515–530.
- Cosmides, L. and J. Tooby (1989). “Evolutionary psychology and the generation of culture, part II: Case study: A computational theory of social exchange”. In: *Ethology and Sociobiology* 10, pp. 51–97.

- Feldman, H. and K. Friston (2010). "Attention, uncertainty, and free-energy". In: *Frontiers in Human Neuroscience* 4, pp. 1–23.
- Fletcher, P. C. and C. D. Frith (2009). "Perceiving is believing: a Bayesian approach to explaining the positive symptoms of schizophrenia". In: *Nature Reviews Neuroscience* 10, pp. 48–58.
- Franklin, D. W. and D. M. Wolpert (2011). "Computational mechanisms of sensor-motor control". In: *Neuron* 72, pp. 425–442.
- Friston, K. (2003). "Learning and inference in the brain". In: *Neural Networks* 16, pp. 1325–1352.
- (2005). "A theory of cortical responses". In: *Philosophical Transactions of the Royal Society of London, Series B* 360, pp. 815–836.
- (2009). "The free-energy principle: a rough guide to the brain?" In: *Trends in Cognitive Sciences* 13, pp. 293–301.
- (2010). "The free-energy principle: A unified brain theory?" In: *Nature Reviews Neuroscience* 11, pp. 127–138.
- (2011). "What is optimal about motor control?" In: *Neuron* 72, pp. 488–498.
- (2013). "Life as we know it". In: *Journal of the Royal Society Interface* 10, p. 20130475.
- Friston, K., J. Daunizeau, J. Kilner and S. J. Kiebel (2010). "Action and behavior: A free-energy formulation". In: *Biological Cybernetics* 102, pp. 227–260.
- Friston, K., J. Kilner and L. Harrison (2006). "A free energy principle for the brain". In: *Journal of Physiology (Paris)* 100, pp. 70–87.
- Friston, K., J. Mattout and J. Kilner (2011). "Action understanding and active inference". In: *Biological Cybernetics* 104, pp. 137–160.
- Friston, K. and K. E. Stephan (2007). "Free-energy and the brain". In: *Synthese* 159, pp. 417–458.
- Friston, K., K. E. Stephan, P. R. Montague and R. J. Dolan (2014). "Computational psychiatry: the brain as a phantastic organ". In: *The Lancet Psychiatry* 1, pp. 148–158.
- Friston, K., C. Thornton and A. Clark (2012). "Free-energy minimization and the dark-room problem". In: *Frontiers in Psychology* 3, pp. 1–7.
- Hohwy, J. (2013). *The Predictive Mind*. Oxford: Oxford University Press.

- Hohwy, J. and J. Michael (2017). “Why should any body have a self?” In: *The Body and the Self, Revisited*. Ed. by F. de Vignemont and A. Alsmith. Cambridge, MA: MIT Press, pp. 363–392.
- Hohwy, J., A. Roepstorff and K. Friston (2008). “Predictive coding explains binocular rivalry: An epistemological review”. In: *Cognition* 108, pp. 687–701.
- Hosoya, T., S. A. Baccus and M. Meister (2005). “Dynamic predictive coding by the retina”. In: *Nature* 436, pp. 71–77.
- Huang, G. T. (May 2008). “Is this a unified theory of the brain?” In: *New Scientist* 2658, pp. 30–33.
- Jehee, J. F. M. and D. H. Ballard (2009). “Predictive feedback can account for biphasic responses in the lateral geniculate nucleus”. In: *PLoS Computational Biology* 5, e1000373.
- Kirchhoff, M. D. and J. Kiverstein (2019). “How to determine the boundaries of the mind: A Markov blanket proposal”. In: *Synthese*. DOI: [10.1007/s11229-019-02370-y](https://doi.org/10.1007/s11229-019-02370-y).
- Koelsch, S., P. Vuust and K. Friston (2019). “Predictive processes and the peculiar case of music”. In: *Trends in Cognitive Sciences* 23, pp. 63–77.
- Kok, P. and F. P. de Lange (2015). “Predictive coding in the sensory cortex”. In: *An Introduction to Model-Based Cognitive Neuroscience*. Ed. by B. U. Forstmann and E.- J. Wagenmakers. New York, NY: Springer, pp. 221–244.
- Kok, P., J. F. M. Jehee and F. P. de Lange (2012). “Less is more: Expectation sharpens representations in the primary visual cortex”. In: *Neuron* 75, pp. 265–270.
- Kording, K. (2007). “Decision theory: What “should” the nervous system do?” In: *Science* 318, pp. 606–610.
- Lupyan, G. (2015). “Cognitive penetrability of perception in the age of prediction: Predictive systems are penetrable systems”. In: *Review of Philosophy and Psychology* 6, pp. 547–569.
- Machery, E. (forthcoming). “Discovery and confirmation in evolutionary psychology”. In: *The Oxford Handbook of Philosophy of Psychology*. Ed. by J. Prinz. Oxford University Press.
- Marr, D. (1982). *Vision*. San Francisco, CA: W. H. Freeman.
- Miller, M. and A. Clark (2018). “Happily entangled: prediction, emotion, and the embodied mind”. In: *Synthese* 195, pp. 2559–2575.

- Muckli, L. (2010). "What are we missing here? Brain imaging evidence for higher cognitive functions in primary visual cortex V1". In: *International Journal of Imaging Systems and Technology* 20, pp. 131–139.
- Mumford, D. (1992). "On the computational architecture of the neocortex: II The role of cortico-cortico loops". In: *Biological Cybernetics* 66, pp. 241–251.
- Murray, S. O., D. Kersten, B. A. Olshausen, P. Schrater and D. L. Woods (2002). "Shape perception reduces activity in human primary visual cortex". In: *Proceedings of the National Academy of Sciences* 99, pp. 15164–15169.
- Olshausen, B. A. and D. J. Field (2005). "How close are we to understanding V1". In: *Neural Computation* 17, pp. 1665–1699.
- Pellicano, E. and D. Burr (2012). "When the world becomes 'too real': a Bayesian explanation of autistic perception". In: *Trends in Cognitive Sciences* 16, pp. 504–510.
- Pickering, M. J. and A. Clark (2014). "Getting ahead: Forward models and their place in cognitive architecture". In: *Trends in Cognitive Sciences* 18, pp. 451–456.
- Ramstead, M. J. D., M. D. Kirchhoff, A. Constant and K. Friston (2021). "Multiscale integration: beyond internalism and externalism". In: *Synthese* 198, S41–S70.
- Rao, R. P. N. and D. H. Ballard (1999). "Predictive coding in the visual cortex: A functional interpretation of some extra-classical receptive-field effects". In: *Nature Neuroscience* 2, pp. 79–87.
- Rao, R. P. N. and T. J. Sejnowski (2002). "Predictive coding, cortical feedback, and spike-timing dependent plasticity". In: *Probabilistic Models of the Brain: Perception and Neural Function*. Ed. by R. P. N. Rao, B. A. Olshausen and M. S. Lewicki. Cambridge, MA: MIT Press, pp. 297–315.
- Schwartenbeck, P., T. FitzGerald, R. J. Dolan and K. Friston (2013). "Exploration, novelty, surprise, and free energy minimization". In: *Frontiers in Psychology* 4, pp. 1–5.
- Schwarz, O. and E. P. Simoncelli (2001). "Natural signal statistics and sensory gain control". In: *Nature Neuroscience* 4, pp. 819–825.
- Seth, A. K. (2013). "Interoceptive inference, emotion, and the embodied self". In: *Trends in Cognitive Sciences* 17, 565–573.
- Shadmehr, R. and J. W. Krakauer (2008). "A computational neuroanatomy for motor control". In: *Experimental Brain Research* 185, pp. 359–381.
- Shagrir, O. (2010). "Marr on computational-level theories". In: *Philosophy of Science* 77, pp. 477–500.

- Shagrir, O. and W. Bechtel (2013). “Marr’s computational level and delineating phenomena”. In: *Integrating Psychology and Neuroscience: Prospects and Problems*. Ed. by D. Kaplan. Oxford: Oxford University Press.
- Spratling, M. W. (2010). “Predictive coding as a model of response properties in cortical area V1”. In: *Journal of Neuroscience* 30, pp. 3531–3543.
- (2017). “A review of predictive coding algorithms”. In: *Brain and Cognition* 112, pp. 92–97.
- Sprevak, M. (2020). “Two kinds of information processing in cognition”. In: *Review of Philosophy and Psychology* 11, pp. 591–611.
- (forthcoming[a]). “Predictive coding I: Introduction”. In: *TBC*.
- (forthcoming[b]). “Predictive coding III: The algorithmic level”. In: *TBC*.
- (forthcoming[c]). “Predictive coding IV: The implementation level”. In: *TBC*.
- (forthcoming[d]). “Predictive coding: Appendix”. In: *TBC*.
- Srinivasan, M. V., S. B. Laughlin and A. Dubs (1982). “Predictive coding: A fresh view of inhibition in the retina”. In: *Proceedings of the Royal Society, Series B* 216, pp. 427–459.
- Wiese, W. (2017). “Action is enabled by systematic misrepresentation”. In: *Erkenntnis* 82, pp. 1233–1252.